

Simulating the Raft Consensus Protocol

Rex Fernando

Dec 17, 2015

Goals

Problem: Comparison of Raft and Paxos

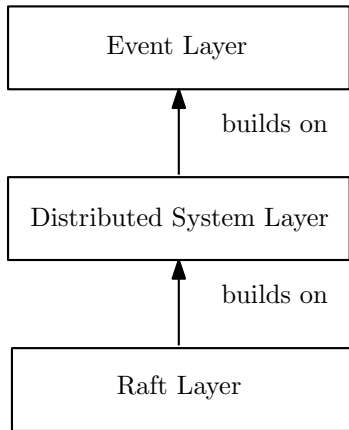
- ▶ Behavior in edge cases
- ▶ Ease of implementing

Proto-Goals

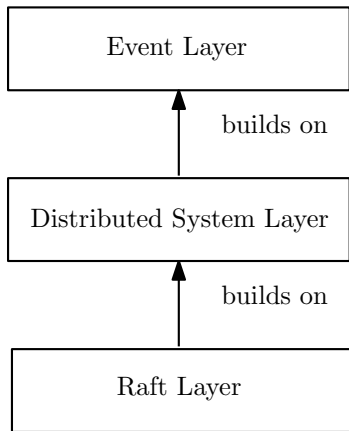
Problem: Build a useful simulation framework

- ▶ Should be able to simulate Raft and Paxos
- ▶ For now, only Raft
- ▶ Goal: generality

Solution/Outline



Solution/Outline



- ▶ Sampling of each piece
- ▶ Example Simulation Runs

Event layer

- ▶ Input: sequence of events, handler, starting state
- ▶ Handler: (State, Event) \rightarrow (Modified state, List of events to emit)
- ▶ Output: sequence of events, states (“filled out”)

Event layer

- ▶ Input: sequence of events, handler, starting state
- ▶ Handler: (State, Event) \rightarrow (Modified state, List of events to emit)
- ▶ Output: sequence of events, states (“filled out”)

Example

Event layer

- ▶ Input: sequence of events, handler, starting state
- ▶ Handler: (State, Event) \rightarrow (Modified state, List of events to emit)
- ▶ Output: sequence of events, states (“filled out”)

Example

- ▶ Input
- ▶ Events: [Switch on at time 0, Switch off at time 1]
- ▶ Starting state: Light is off
- ▶ Handler:
 - ▶ If switch is switched on at time t , emit light-on event at $t+0.01$.
 - ▶ If switch is switched off at time t , emit light-off event at $t+0.01$.

Event layer

- ▶ Input: sequence of events, handler, starting state
- ▶ Handler: (State, Event) \rightarrow (Modified state, List of events to emit)
- ▶ Output: sequence of events, states (“filled out”)

Example

- ▶ Input
- ▶ Events: [Switch on at time 0, Switch off at time 1]
- ▶ Starting state: Light is off
- ▶ Handler:
 - ▶ If switch is switched on at time t , emit light-on event at $t+0.01$.
 - ▶ If switch is switched off at time t , emit light-off event at $t+0.01$.
- ▶ Output
- ▶ Events: [Switch on at time 0, Light on at time 0.01, Switch off at time 1, Light off at time 1.01]

Dist. System Layer

- ▶ System now consists of n machines
- ▶ Input:
- ▶ Sequence of events
- ▶ Handlers and starting states *for each machine in the system*
- ▶ *Global system behavior description*
 - ▶ Network
 - ▶ Crashes
 - ▶ Input into the system
- ▶ Output: sequence of events, states

Candidates (§5.2):

- On conversion to candidate, start election:
 - Increment currentTerm
 - Vote for self
 - Reset election timer
 - Send RequestVote RPCs to all other servers
- If votes received from majority of servers: become leader
- If AppendEntries RPC received from new leader: convert to follower
- If election timeout elapses: start new election



Raft Layer

```
1  -- Response handler
2  raftHandler (ReqVR term voteGranted) = do
3      state <- get
4      if term > (currentTerm state) then do
5          becomeFollower term (-1)
6      else if voteGranted then do
7          put state { votesForMe = (votesForMe state)+1 }
8          state <- get
9          if (votesForMe state) >= quorum then
10             becomeLeader
11             else return ()
12             else return ()
```

Crash after Elect

```
1  global (Event time (Receive _ leader _
2      (AppE _ _ _ _ _))) ms = do
3      alreadyCrashed <- get
4      if not alreadyCrashed then do
5          crash leader
6          put True
7      else return ()
8      sendAllMessages time ms
9
10 sendAllMessages time ms =
11     mapM_ (flip send (time+delay)) ms
```


Crash after Elect




Crash after Elect 2

```
1  global (Event time (Receive _ leader _
2              (AppE _ _ _ _ _))) ms = do
3      alreadyCrashed <- get
4      if not alreadyCrashed then do
5          crash leader
6          --put True
7          else return ()
8      sendAllMessages time ms
9
10 sendAllMessages time ms =
11     mapM_ (flip send (time+delay)) ms
```

Crash after Elect 2



Crash during log write



Summary

Simulation framework

- ▶ Describe a distributed system (in this case Raft)
- ▶ Give custom “global behavior”
- ▶ Visualize results